

Dare 2023 – Summer School



NOVA SCHOOL OF
SCIENCE & TECHNOLOGY

André Costa and David Antunes

TLA+: 2PC and 3PC Spec

Introduction

- Two Phase Commit
- Leslie Lamport's Spec

Introduction

- Two Phase Commit
- Leslie Lamport's Spec

Specification of Atomic Commit

Specification of Atomic Commit

Model:

Crash Recovery
Synchronous System

Specification of Atomic Commit

Model: Crash Recovery

Agreement: Any two processes that decide, decide the same value.

Specification of Atomic Commit

Model: Crash Recovery

Agreement: Any two processes that decide, decide the same value.

Validity (1): If some process starts with the value “no” then “abort” is the only possible decision.

Specification of Atomic Commit

Model: Crash Recovery

Agreement: Any two processes that decide, decide the same value.

Validity (1): If some process starts with the value “no” then “abort” is the only possible decision.

Validity (2): If all processes start with value “yes” and none fails, then “commit” is the only possible decision.

Specification of Atomic Commit

Model: Crash Recovery

Agreement: Any two processes that decide, decide the same value.

Validity (1): If some process starts with the value “no” then “abort” is the only possible decision.

Validity (2): If all processes start with value “yes” and none fails, then “commit” is the only possible decision.

Termination: If eventually all processes recover from all faults, then, eventually all processes decide.

Two Phase Commit

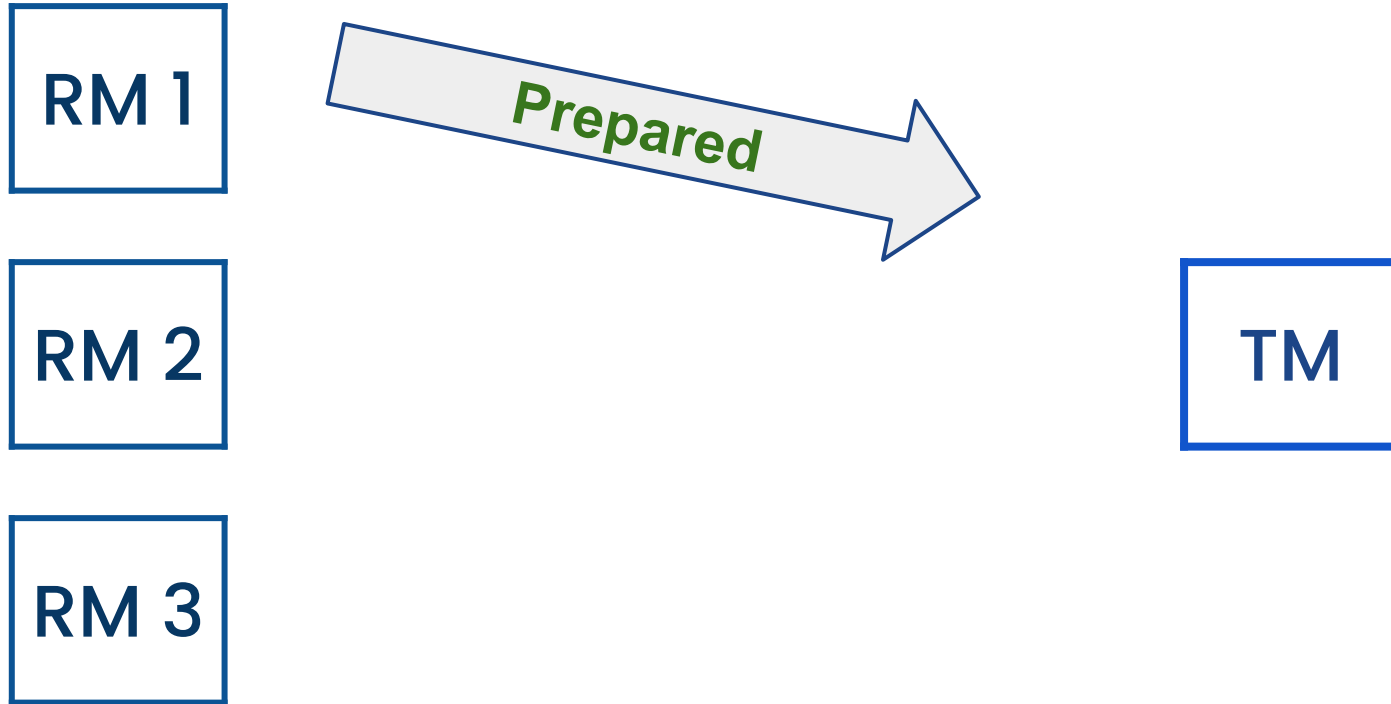
RM 1

RM 2


RM 3

TM

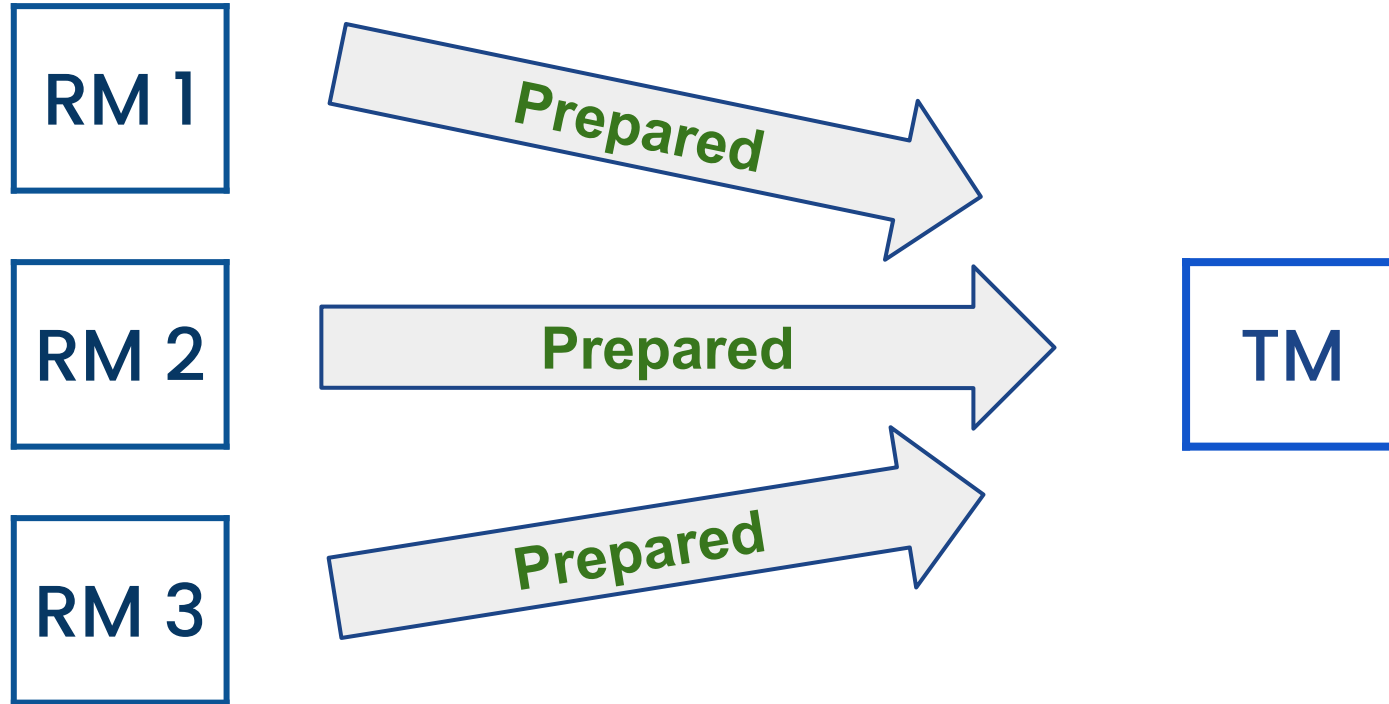
Two Phase Commit




Two Phase Commit



Two Phase Commit



Two Phase Commit



Two Phase Commit

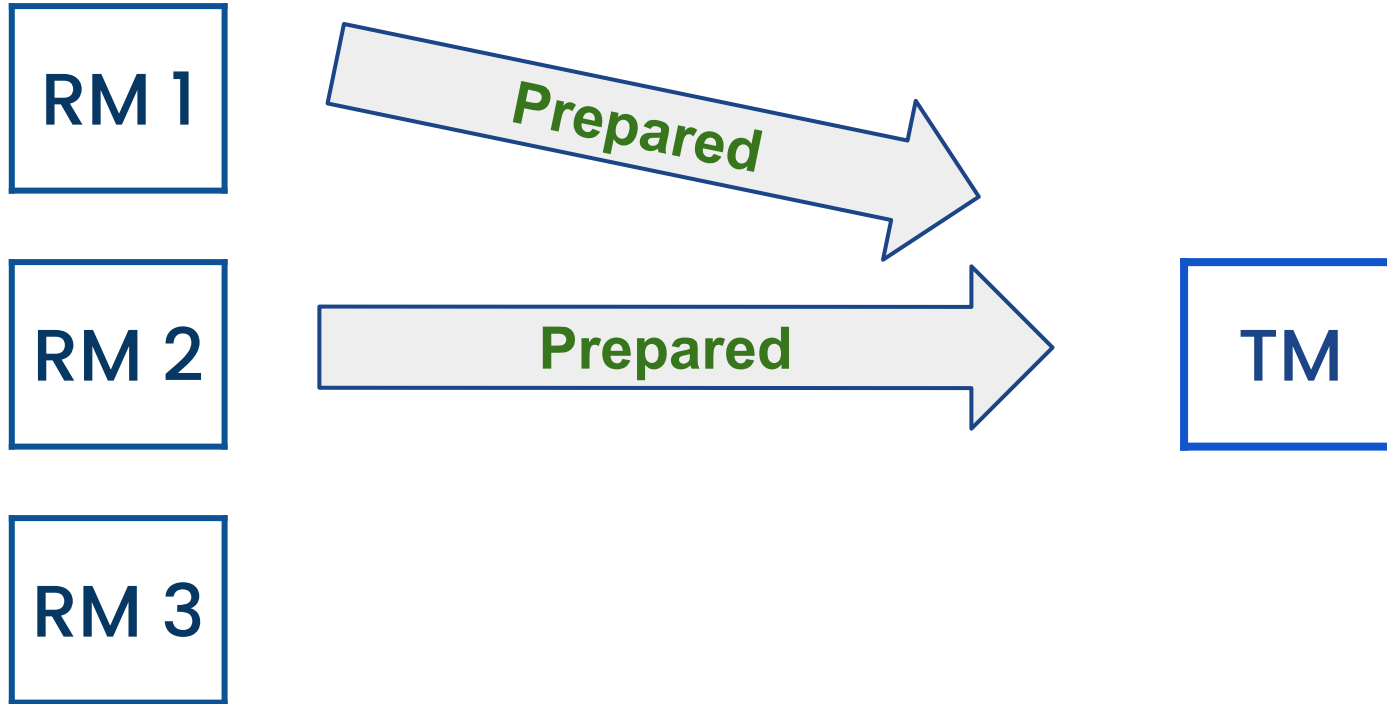
RM 1

RM 2

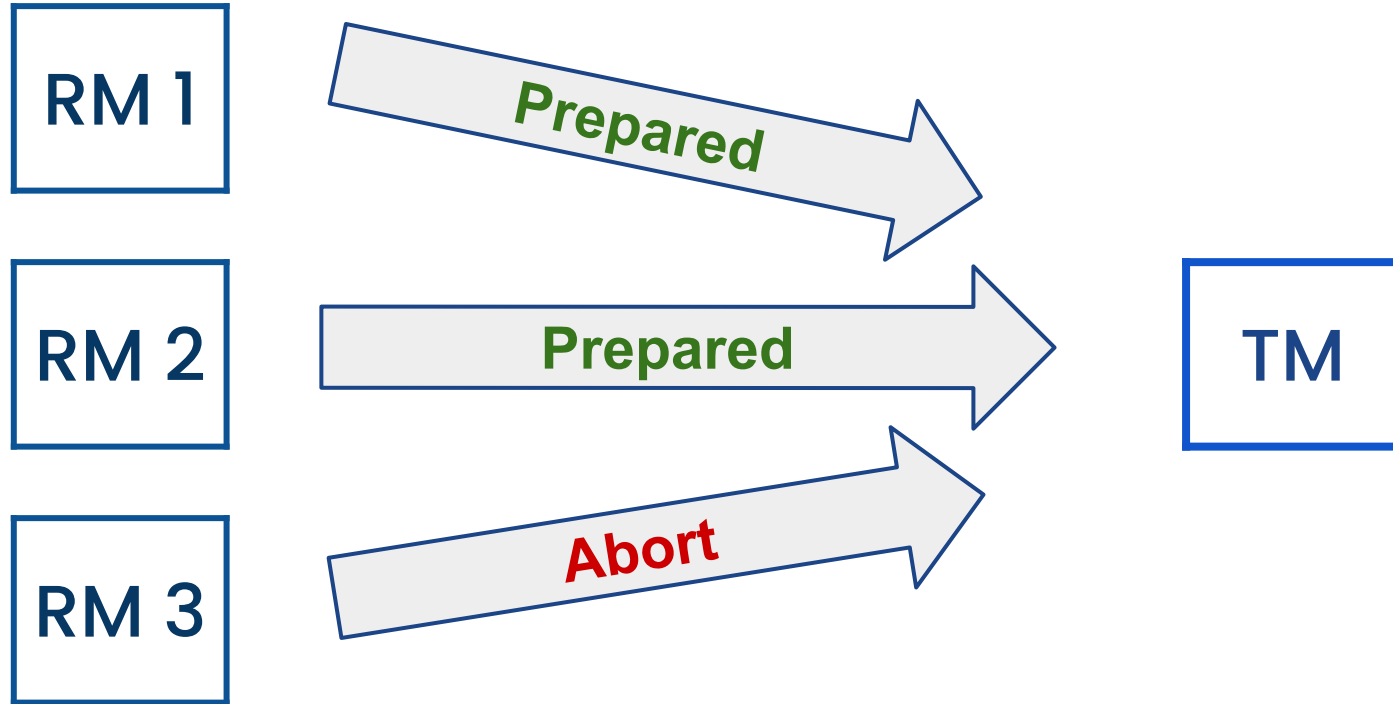
RM 3

TM


Two Phase Commit



Two Phase Commit



Two Phase Commit



Introduction

- Two Phase Commit
- Leslie Lamport's Spec

Introduction

- Two Phase Commit
- Leslie Lamport's Spec

Leslie Lamport's Specification

Leslie Lamport's Specification

CONSTANT RM * The set of resource managers

Leslie Lamport's Specification

CONSTANT RM * The set of resource managers

VARIABLES

Leslie Lamport's Specification

CONSTANT RM * The set of resource managers

VARIABLES

rmState * State of each resource manager.

Leslie Lamport's Specification

CONSTANT RM * The set of resource managers

VARIABLES

rmState * State of each resource manager.

tmState * The state of the transaction manager.

Leslie Lamport's Specification

CONSTANT RM * The set of resource managers

VARIABLES

rmState * State of each resource manager.

tmState * The state of the transaction manager.

tmPrepared * The set of RMs from which the TM has received prepared messages

Leslie Lamport's Specification

CONSTANT RM \ * The set of resource managers

VARIABLES

rmState \ * State of each resource manager.

tmState \ * The state of the transaction manager.

tmPrepared \ * The set of RMs from which the TM has received prepared messages

msgs \ * The set of all messages ever sent

Leslie Lamport's Specification

Message ==

Leslie Lamport's Specification

Message ==

$[type : \{\text{"Prepared"}\}, rm : RM] \cup [type : \{\text{"Commit"}, \text{"Abort"}\}]$

Leslie Lamport's Specification

Message ==

$[type : \{\text{"Prepared"}\}, rm : RM] \cup [type : \{\text{"Commit"}, \text{"Abort"}\}]$

TPTYPEOK ==

Leslie Lamport's Specification

Message ==

$[type : \{ "Prepared" \}, rm : RM] \cup [type : \{ "Commit", "Abort" \}]$

TPTYPEOK ==

$\wedge rmState \in [RM \rightarrow \{ "working", "prepared", "committed", "aborted" \}]$

Leslie Lamport's Specification

Message ==

$[type : \{ "Prepared" \}, rm : RM] \cup [type : \{ "Commit", "Abort" \}]$

TPTYPEOK ==

$\wedge rmState \in [RM \rightarrow \{ "working", "prepared", "committed", "aborted" \}]$
 $\wedge tmState \in \{ "init", "committed", "aborted" \}$

Leslie Lamport's Specification

Message ==

[type : {"Prepared"}, rm : RM] \cup [type : {"Commit", "Abort"}]

TPTYPEOK ==

\wedge rmState \in [RM \rightarrow {"working", "prepared", "committed", "aborted"}]

\wedge tmState \in {"init", "committed", "aborted"}

\wedge tmPrepared \subseteq RM

Leslie Lamport's Specification

Message ==

[type : {"Prepared"}, rm : RM] \cup [type : {"Commit", "Abort"}]

TPTYPEOK ==

\wedge rmState \in [RM \rightarrow {"working", "prepared", "committed", "aborted"}]

\wedge tmState \in {"init", "committed", "aborted"}

\wedge tmPrepared \subseteq RM

\wedge msgs \subseteq Message

Leslie Lamport's Specification

(Actions that lead to Commit)

Leslie Lamport's Specification

(Actions that lead to Commit)

\ * Resource Manager sends prepare message

RMPPrepare(rm) ==

Leslie Lamport's Specification

(Actions that lead to Commit)

* Resource Manager sends prepare message

RMPPrepare(rm) ==

/\ rmState[rm] = "working"

Leslie Lamport's Specification

(Actions that lead to Commit)

* Resource Manager sends prepare message

RMPPrepare(rm) ==

/\ rmState[rm] = "working"

/\ rmState' = [rmState EXCEPT ![rm] = "prepared"]

Leslie Lamport's Specification

(Actions that lead to Commit)

* Resource Manager sends prepare message

RMPPrepare(rm) ==

/\ rmState[rm] = "working"

/\ rmState' = [rmState EXCEPT ![rm] = "prepared"]

/\ msgs' = msgs \cup {[type |-> "Prepared", rm |-> rm]}

Leslie Lamport's Specification

(Actions that lead to Commit)

* Resource Manager sends prepare message

RMPPrepare(rm) ==

/\ rmState[rm] = "working"

/\ rmState' = [rmState EXCEPT ![rm] = "prepared"]

/\ msgs' = msgs \cup {[type |-> "Prepared", rm |-> rm]}

* Transaction Manager receives prepare message

TMRcvPrepared(rm) ==

Leslie Lamport's Specification

(Actions that lead to Commit)

* Resource Manager sends prepare message

RMPPrepare(rm) ==

/\ rmState[rm] = "working"

/\ rmState' = [rmState EXCEPT ![rm] = "prepared"]

/\ msgs' = msgs \cup {[type |-> "Prepared", rm |-> rm]}

* Transaction Manager receives prepare message

TMRcvPrepared(rm) ==

/\ tmState = "init"

Leslie Lamport's Specification

(Actions that lead to Commit)

* Resource Manager sends prepare message

RMPPrepare(rm) ==

/\ rmState[rm] = "working"

/\ rmState' = [rmState EXCEPT ![rm] = "prepared"]

/\ msgs' = msgs \cup {[type |-> "Prepared", rm |-> rm]}

* Transaction Manager receives prepare message

TMRcvPrepared(rm) ==

/\ tmState = "init"

/\ [type |-> "Prepared", rm |-> rm] \in msgs

Leslie Lamport's Specification

(Actions that lead to Commit)

* Resource Manager sends prepare message

RMPPrepare(rm) ==

/\ rmState[rm] = "working"

/\ rmState' = [rmState EXCEPT ![rm] = "prepared"]

/\ msgs' = msgs \cup {[type |-> "Prepared", rm |-> rm]}

* Transaction Manager receives prepare message

TMRcvPrepared(rm) ==

/\ tmState = "init"

/\ [type |-> "Prepared", rm |-> rm] \in msgs

/\ tmPrepared' = tmPrepared \cup {rm}

Leslie Lamport's Specification

(Actions that lead to Commit)

Leslie Lamport's Specification

(Actions that lead to Commit)

* Transaction Manager sends **Commit** message

TMCommit ==

Leslie Lamport's Specification

(Actions that lead to Commit)

* Transaction Manager sends **Commit** message

TMCommit ==

/\ tmState = "init"

Leslie Lamport's Specification

(Actions that lead to Commit)

* Transaction Manager sends **Commit** message

TMCommit ==

/\ tmState = "init"

/\ tmPrepared = RM

Leslie Lamport's Specification

(Actions that lead to Commit)

* Transaction Manager sends **Commit** message

TMCommit ==

/\ tmState = "init"

/\ tmPrepared = RM

/\ tmState' = "committed"

Leslie Lamport's Specification

(Actions that lead to Commit)

* Transaction Manager sends **Commit** message

TMCommit ==

/\ tmState = "init"

/\ tmPrepared = RM

/\ tmState' = "committed"

/\ msgs' = msgs \cup {[type |-> "Commit"]}

Leslie Lamport's Specification

(Actions that lead to Commit)

* Transaction Manager sends **Commit** message

TMCommit ==

/\ tmState = "init"

/\ tmPrepared = RM

/\ tmState' = "committed"

/\ msgs' = msgs \cup {[type |-> "Commit"]}

* Resource Manager receives **Commit** message

RMRcvCommitMsg(rm) ==

Leslie Lamport's Specification

(Actions that lead to Commit)

* Transaction Manager sends **Commit** message

TMCommit ==

/\ tmState = "init"

/\ tmPrepared = RM

/\ tmState' = "committed"

/\ msgs' = msgs \cup {[type |-> "Commit"]}

* Resource Manager receives **Commit** message

RMRcvCommitMsg(rm) ==

/\ [type |-> "Commit"] \in msgs

Leslie Lamport's Specification

(Actions that lead to Commit)

* Transaction Manager sends **Commit** message

TMCommit ==

/\ tmState = "init"

/\ tmPrepared = RM

/\ tmState' = "committed"

/\ msgs' = msgs \cup {[type |-> "Commit"]}

* Resource Manager receives **Commit** message

RMRCvCommitMsg(rm) ==

/\ [type |-> "Commit"] \in msgs

/\ rmState' = [rmState EXCEPT ![rm] = "committed"]

Leslie Lamport's Specification

Leslie Lamport's Specification

**Transaction Manager can decide to abort
(even when all RMs want to commit)**

Leslie Lamport's Specification

Transaction Manager can decide to abort
(even when all RMs want to commit)

* The TM spontaneously aborts the transaction.

TMAbort ==

/\ tmState = "init"

Leslie Lamport's Specification

Transaction Manager can decide to abort
(even when all RMs want to commit)

* The TM spontaneously aborts the transaction.

TMAbort ==

/\ tmState = "init"

/\ tmState' = "aborted"

Leslie Lamport's Specification

Transaction Manager can decide to abort
(even when all RMs want to commit)

* The TM spontaneously aborts the transaction.

TMAbort ==

/\ tmState = "init"

/\ tmState' = "aborted"

/\ msgs' = msgs \cup {[type |-> "Abort"]}

Our 2PC Specification

Our 2PC Specification

Contributions:

Our 2PC Specification

Contributions:

- The Transaction Manager can no longer spontaneously abort transactions.

Our 2PC Specification

Contributions:

- The Transaction Manager can no longer spontaneously abort transactions.
- Resource Managers can crash.

Our 2PC Specification

Contributions:

- The Transaction Manager can no longer spontaneously abort transactions.
- Resource Managers can crash.
- Every Atomic Commit Property is checked.

Our 2PC Specification

**The Transaction Manager can no longer
spontaneously abort transactions:**

Our 2PC Specification

The Transaction Manager can no longer
spontaneously abort transactions:

* A RM decides to abort.

RMChooseToAbort(rm) ==
/\ rmState[rm] = "working"

Our 2PC Specification

The Transaction Manager can no longer
spontaneously abort transactions:

* A RM decides to abort.

RMChooseToAbort(rm) ==

/\ rmState[rm] = "working"

/\ rmState' = [rmState EXCEPT ![rm] = "aborted"]

/\ msgs' = msgs \cup {[type |-> "Aborted", rm |-> rm]}

Our 2PC Specification

The Transaction Manager can no longer spontaneously abort transactions:

* A RM decides to abort.

RMChooseToAbort(rm) ==

 /\ rmState[rm] = "working"

 /\ rmState' = [rmState EXCEPT ![rm] = "aborted"]

 /\ msgs' = msgs \cup {[type |-> "Aborted", rm |-> rm]}

* The TM aborts the transaction, as requested by some RM.

TMAbort ==

Our 2PC Specification

The Transaction Manager can no longer spontaneously abort transactions:

* A RM decides to abort.

RMChooseToAbort(rm) ==

/\ rmState[rm] = "working"

/\ rmState' = [rmState EXCEPT ![rm] = "aborted"]

/\ msgs' = msgs \cup {[type |-> "Aborted", rm |-> rm]}

* The TM aborts the transaction, as requested by some RM.

TMAbort ==

/\ tmState = "init"

/\ [type |-> "Aborted", rm |-> rm] \in msgs

Our 2PC Specification

The Transaction Manager can no longer spontaneously abort transactions:

* A RM decides to abort.

RMChooseToAbort(rm) ==

/\ rmState[rm] = "working"

/\ rmState' = [rmState EXCEPT ![rm] = "aborted"]

/\ msgs' = msgs \cup {[type |-> "Aborted", rm |-> rm]}

* The TM aborts the transaction, as requested by some RM.

TMAbort ==

/\ tmState = "init"

/\ [type |-> "Aborted", rm |-> rm] \in msgs

/\ tmState' = "aborted"

/\ msgs' = msgs \cup {[type |-> "Abort"]}

Our 2PC Specification

Resource Managers can crash:

Our 2PC Specification

Resource Managers can crash:

* A RM crashes.

RMCrash(rm) ==

Our 2PC Specification

Resource Managers can crash:

* A RM crashes.

RMCrash(rm) ==
 /\ rmState[rm] /= "crashed"

Our 2PC Specification

Resource Managers can crash:

* A RM crashes.

RMCrash(rm) ==

/\ rmState[rm] /= "crashed"

/\ rmPrevState' = [rmPrevState EXCEPT ![rm] = rmState[rm]]

Our 2PC Specification

Resource Managers can crash:

* A RM crashes.

RMCrash(rm) ==

/\ rmState[rm] /= "crashed"

/\ rmPrevState' = [rmPrevState EXCEPT ![rm] = rmState[rm]]

/\ rmState' = [rmState EXCEPT ![rm] = "crashed"]

Our 2PC Specification

Resource Managers can crash:

* A RM crashes.

RMCrash(rm) ==

 /\ rmState[rm] /= "crashed"

 /\ rmPrevState' = [rmPrevState EXCEPT ![rm] = rmState[rm]]

 /\ rmState' = [rmState EXCEPT ![rm] = "crashed"]

* A RM recovers from a crash.

RMRecover(rm) ==

Our 2PC Specification

Resource Managers can crash:

* A RM crashes.

RMCrash(rm) ==

/\ rmState[rm] /= "crashed"

/\ rmPrevState' = [rmPrevState EXCEPT ![rm] = rmState[rm]]

/\ rmState' = [rmState EXCEPT ![rm] = "crashed"]

* A RM recovers from a crash.

RMRecover(rm) ==

/\ rmState[rm] = "crashed"

Our 2PC Specification

Resource Managers can crash:

* A RM crashes.

RMCrash(rm) ==

/\ rmState[rm] /= "crashed"

/\ rmPrevState' = [rmPrevState EXCEPT ![rm] = rmState[rm]]

/\ rmState' = [rmState EXCEPT ![rm] = "crashed"]

* A RM recovers from a crash.

RMRecover(rm) ==

/\ rmState[rm] = "crashed"

/\ rmState' = [rmState EXCEPT ![rm] = rmPrevState[rm]]

Our 2PC Specification

Resource Managers can crash:

* A RM crashes.

RMCrash(rm) ==

/\ rmState[rm] /= "crashed"

/\ rmPrevState' = [rmPrevState EXCEPT ![rm] = rmState[rm]]

/\ rmState' = [rmState EXCEPT ![rm] = "crashed"]

* A RM recovers from a crash.

RMRecover(rm) ==

/\ rmState[rm] = "crashed"

/\ rmState' = [rmState EXCEPT ![rm] = rmPrevState[rm]]

/\ rmPrevState' = [rmPrevState EXCEPT ![rm] = "crashed"]

Our 2PC Specification

Every Atomic Commit Property is checked:

Our 2PC Specification

Every Atomic Commit Property is checked:

Agreement: Any two processes that decide, decide the same value

Our 2PC Specification

Every Atomic Commit Property is checked:

Agreement: Any two processes that decide, decide the same value

TPAgreement ==

$\nexists A \text{ rm1, rm2 } \nexists \text{ in RM :}$

$\sim \wedge \text{rmState[rm1]} = \text{"aborted"}$

$\wedge \text{rmState[rm2]} = \text{"committed"}$

Our 2PC Specification

Every Atomic Commit Property is checked:

Agreement: Any two processes that decide, decide the same value

TPAgreement ==

$\nexists \text{rm1, rm2} \in \text{RM} :$

$\sim \wedge \text{rmState}[\text{rm1}] = \text{"aborted"}$

$\wedge \text{rmState}[\text{rm2}] = \text{"committed"}$

Translation: No two processes can ever be in different decision states (i.e., decide different things).

Our 2PC Specification

Every Atomic Commit Property is checked:

Our 2PC Specification

Every Atomic Commit Property is checked:

Validity (1): If some process starts with the value “no” then “abort” is the only possible decision

Our 2PC Specification

Every Atomic Commit Property is checked:

Validity (1): If some process starts with the value “no” then “abort” is the only possible decision

TPValidity1 ==

$(\exists rm \in RM : rmState[rm] = \text{"aborted"}) \sim \rightarrow (tmState = \text{"aborted"})$

Our 2PC Specification

Every Atomic Commit Property is checked:

Validity (1): If some process starts with the value “no” then “abort” is the only possible decision

TPValidity1 ==

$(\exists rm \in RM : rmState[rm] = \text{"aborted"}) \sim \rightarrow (tmState = \text{"aborted"})$

Translation: If any process is in aborted state (decided abort), then TM will eventually decide abort.

Our 2PC Specification

Every Atomic Commit Property is checked:

Our 2PC Specification

Every Atomic Commit Property is checked:

Validity (2): If all processes start with value “yes” and none fails, then “commit” is the only possible decision

Our 2PC Specification

Every Atomic Commit Property is checked:

Validity (2): If all processes start with value “yes” and none fails, then “commit” is the only possible decision

TPValidity2 ==

$\sim (\exists rm \in RM : rmState[rm] \neq \text{"prepared"}) \sim \rightarrow (tmState = \text{"committed"})$

Our 2PC Specification

Every Atomic Commit Property is checked:

Validity (2): If all processes start with value “yes” and none fails, then “commit” is the only possible decision

TPValidity2 ==

$\sim (\exists rm \in RM : rmState[rm] \neq \text{"prepared"}) \sim \rightarrow (tmState = \text{"committed"})$

Translation: If all processes are prepared (which also means it is not crashed), then eventually TM will decide commit.

Our 2PC Specification

Every Atomic Commit Property is checked:

Our 2PC Specification

Every Atomic Commit Property is checked:

Termination: If eventually all processes recover from all faults, then, eventually all processes decide

Our 2PC Specification

Every Atomic Commit Property is checked:

Termination: If eventually all processes recover from all faults, then, eventually all processes decide

TPTermination ==

$$\begin{aligned} & (\sim (\exists rm \in RM : rmState[rm] = "crashed")) \sim \rightarrow \\ & \quad (\sim (\exists rm \in RM : rmState[rm] \neq "aborted" \wedge rmPrevState[rm] \neq "aborted")) \\ & \quad \wedge (\sim (\exists rm \in RM : rmState[rm] \neq "committed" \wedge rmPrevState[rm] \neq "committed")) \end{aligned}$$

Our 2PC Specification

Every Atomic Commit Property is checked:

Termination: If eventually all processes recover from all faults, then, eventually all processes decide

TPTermination ==

$$\begin{aligned} & (\sim (\exists rm \in RM : rmState[rm] = "crashed")) \sim \rightarrow \\ & \quad (\sim (\exists rm \in RM : rmState[rm] \neq "aborted" \wedge rmPrevState[rm] \neq "aborted")) \\ & \quad \vee (\sim (\exists rm \in RM : rmState[rm] \neq "committed" \wedge rmPrevState[rm] \neq "committed")) \end{aligned}$$

Translation: If eventually no process is crashed, then eventually all will decide (the same thing). A process has decided if: it is in a decided state OR it was in a decided state (and is now crashed).

Our 2PC Specification

TLC reports no errors!

Time	Diameter	States Found	Distinct States	Queue Size
00:04:22	17	61 396	9 756	0

Our Specification

- Two Phase Commit
- Three Phase Commit

Our Specification

- Two Phase Commit
- Three Phase Commit

Three-Phase Commit

- Converts 2PC into a non blocking protocol
- Timeouts
- Adds Pre-commit round


Three-Phase Commit




Three-Phase Commit




Three-Phase Commit




Three-Phase Commit



Three-Phase Commit



Three-Phase Commit




Three-Phase Commit



Three-Phase Commit



Three-Phase Commit




Three-Phase Commit




Three-Phase Commit




Three-Phase Commit



Three-Phase Commit



Three-Phase Commit




Three-Phase Commit




Three-Phase Commit




Three-Phase Commit




Three-Phase Commit



Three-Phase Commit



Three-Phase Commit



Three-Phase Commit

- Leslie Lamport 2 Phase Commit Specification
- Timeouts
- Pre-commit round

Three-Phase Commit

- Leslie Lamport 2 Phase Commit Specification
- Timeouts
- Pre-commit round

Three-Phase Commit

- Leslie Lamport 2 Phase Commit Specification
- Timeouts
- Pre-commit round

Three-Phase Commit

- Leslie Lamport 2 Phase Commit Specification
- Timeouts
- Pre-commit round

Specification of Atomic Commit

3 Phase Commit Specification:

Three-Phase Commit

CONSTANT RM * The set of resource managers

VARIABLE

rmState, * State of each resource manager

tmState, * State of the transaction manager

rmPrepared, * The set of RMs from which the TM has received prepared messages

rmPrecommitted, * The set of RMs from which the TM has received pre-committed messages

msgs, * Set of all messages

timeout * Timeout flag

Three-Phase Commit

\ * All possible states the RMs and TM can assume

_WORKING == "working"

_PREPARED == "prepared"

_PREPARE == "prepare"

_PRECOMMIT == "pre-commit"

_PRECOMMITTED == "pre-committed"

_COMMIT == "commit"

_COMMITTED == "committed"

_ABORT == "abort"

_ABORTED == "aborted"

_INIT == "init"

Three-Phase Commit

* The set of all possible messages.

Message == [type: { _PREPARE, _PRECOMMIT, _COMMIT, _ABORT }] \cup
[type: { _PREPARED, _PRECOMMITTED, _ABORTED }, rm: RM]

Three-Phase Commit

* The set of all possible messages.

```
Message == [ type: { _PREPARE, _PRECOMMIT, _COMMIT, _ABORT } ] \cup  
            [ type: { _PREPARED, _PRECOMMITTED, _ABORTED }, rm: RM ]
```

* Initial State of the Spec

```
Init ==  
    /\ rmState = [ r \in RM |-> _WORKING ]  
    /\ tmState = _INIT  
    /\ rmPrepared = {}  
    /\ rmPrecommitted = {}  
    /\ msgs = {}  
    /\ timeout = "off"
```

Three-Phase Commit

\ * TM receives a prepare message from r

TMRCvPrepare(r) ==

 /\ timeout = "off"

 /\ tmState = _INIT

 /\ [type |-> _PREPARED, rm |-> r] \in msgs

 /\ r \notin rmPrepared

 /\ rmPrepared' = rmPrepared \cup {r}

 /\ **UNCHANGED**<<tmState, rmState, rmPrecommitted, msgs, timeout>>

Three-Phase Commit

\ * TM receives a prepare message from r

TMrcvPrepare(r) ==

 /\ timeout = "off"

 /\ tmState = _INIT

 /\ [type |-> _PREPARED, rm |-> r] \in msgs

 /\ r \notin rmPrepared

 /\ rmPrepared' = rmPrepared \cup {r}

 /\ **UNCHANGED**<<tmState, rmState, rmPrecommitted, msgs, timeout>>

\ * RM decides to go into prepare state

RMPrepare(r) ==

 /\ timeout = "off"

 /\ rmState[r] = _WORKING

 /\ rmState' = [rmState EXCEPT ![r] = _PREPARED]

 /\ msgs' = msgs \cup { [type |-> _PREPARED, rm |-> r] }

 /\ **UNCHANGED**<<tmState, rmPrepared, rmPrecommitted, timeout>>

Three-Phase Commit

* TM sends pre-commit message after receiving prepared message from RMs

TMPrecommit ==

 /\ timeout = "off"

 /\ tmState = _INIT

 /\ rmPrepared = RM

 /\ rmPrecommitted = {}

 /\ [type |-> _PRECOMMIT] \notin msgs

 /\ msgs' = msgs \cup {[type |-> _PRECOMMIT]}

 /\ **UNCHANGED**<<tmState, rmState, rmPrepared, rmPrecommitted, timeout>>

Three-Phase Commit

```
\ * TM receives a precommit message from r
TMRcvPrecommit(r) ==
  /\ timeout = "off"
  /\ tmState = _INIT
  /\ [ type |-> _PRECOMMITTED, rm |-> r ] \in msgs
  /\ r \notin rmPrecommitted
  /\ rmPrecommitted' = rmPrecommitted \cup {r}
  /\ UNCHANGED<<tmState, rmState, rmPrepared, msgs, timeout>>
```

Three-Phase Commit

* TM receives a precommit message from r

TMRcvPrecommit(r) ==

 /\ timeout = "off"

 /\ tmState = _INIT

 /\ [type |-> _PRECOMMITTED, rm |-> r] \in msgs

 /\ r \notin rmPrecommitted

 /\ rmPrecommitted' = rmPrecommitted \cup {r}

 /\ **UNCHANGED**<<tmState, rmState, rmPrepared, msgs, timeout>>

* RM decides to pre-commit after seeing a pre-commit message from TM

RMPrecommit(r) ==

 /\ timeout = "off"

 /\ rmState[r] = _PREPARED

 /\ [type |-> _PRECOMMIT] \in msgs

 /\ rmState' = [rmState EXCEPT ![r] = _PRECOMMITTED]

 /\ msgs' = msgs \cup { [type |-> _PRECOMMITTED, rm |-> r] }

 /\ **UNCHANGED**<<tmState, rmPrepared, rmPrecommitted, timeout>>

Three-Phase Commit

\ * TM changes to the commit state and sends message to commit

TMCommit ==

 /\ timeout = "off"

 /\ tmState = _INIT

 /\ rmPrecommitted = RM

 /\ tmState' = _COMMITTED

 /\ msgs' = msgs \cup {[type |-> _COMMIT]}

 /\ **UNCHANGED**<<rmState, rmPrepared, rmPrecommitted, timeout>>

Three-Phase Commit

* TM changes to the commit state and sends message to commit

TMCommit ==

```
/\ timeout = "off"  
/\ tmState = _INIT  
/\ rmPrecommitted = RM  
/\ tmState' = _COMMITTED  
/\ msgs' = msgs \cup {[ type |-> _COMMIT ]}  
/\ UNCHANGED<<rmState, rmPrepared, rmPrecommitted, timeout>>
```

* RM receives commit message and changes state to commit

RMCommit(r) ==

```
/\ timeout = "off"  
/\ rmState[r] = _PRECOMMITTED  
/\ [ type |-> _COMMIT ] \in msgs  
/\ rmState' = [ rmState EXCEPT ![r] = _COMMITTED ]  
/\ UNCHANGED<<tmState, rmPrepared, rmPrecommitted, msgs, timeout>>
```

Three-Phase Commit

* TM decides to abort the transaction

TMAbort ==

/\ timeout = "off"

/\ rmPrecommitted = {}

/\ ~ ([type |-> _PRECOMMIT] \in msgs \/ [type |-> _COMMIT] \in msgs)

/\ tmState = _INIT

/\ tmState' = _ABORTED

/\ msgs' = msgs \cup {[type |-> _ABORT]}

/\ **UNCHANGED** <<rmState, rmPrepared, rmPrecommitted, timeout>>

Three-Phase Commit

\ * TM decides to abort the transaction

TMAbort ==

```
/\ timeout = "off"
/\ rmPrecommitted = {}
/\ ~ ([ type |-> _PRECOMMIT ] \in msgs \/ [ type |-> _COMMIT ] \in msgs)
/\ tmState = _INIT
/\ tmState' = _ABORTED
/\ msgs' = msgs \cup {[type |-> _ABORT]}
/\ UNCHANGED <<rmState, rmPrepared, rmPrecommitted, timeout>>
```

\ * TM decides to abort the transaction after seeing an aborted message from an RM and sends abort message

TMRcvAbort(r) ==

```
/\ timeout = "off"
/\ tmState = _INIT
/\ [type |-> _ABORTED, rm |-> r] \in msgs
/\ rmPrecommitted = {}
/\ tmState' = _ABORTED
/\ msgs' = msgs \cup {[type |-> _ABORT]}
/\ UNCHANGED <<rmState, rmPrepared, rmPrecommitted, timeout>>
```

Three-Phase Commit

\ * RM Chooses to abort during the prepare phase

RMChooseToAbort(r) ==

/\ timeout = "off"

/\ rmState[r] = _WORKING

/\ rmState' = [rmState EXCEPT ![r] = _ABORTED]

/\ r \notin rmPrepared

/\ msgs' = msgs \cup {[type |-> _ABORTED, rm |-> r]}

/\ **UNCHANGED** <<tmState, rmPrepared, rmPrecommitted, timeout>>

Three-Phase Commit

\ * RM Chooses to abort during the prepare phase

RMChooseToAbort(r) ==

/\ timeout = "off"

/\ rmState[r] = _WORKING

/\ rmState' = [rmState EXCEPT ![r] = _ABORTED]

/\ r \notin rmPrepared

/\ msgs' = msgs \cup {[type |-> _ABORTED, rm |-> r]}

/\ **UNCHANGED** <<tmState, rmPrepared, rmPrecommitted, timeout>>

\ * RM decides to abort after seeing the abort message from TM

RMRCvAbortMsg(r) ==

/\ timeout = "off"

/\ rmState[r] = _PREPARED /\ rmState[r] = _WORKING

/\ [type |-> _ABORT] \in msgs

/\ rmState' = [rmState EXCEPT ![r] = _ABORTED]

/\ **UNCHANGED** <<tmState, rmPrepared, rmPrecommitted, msgs, timeout>>

Three-Phase Commit

Timeout ==

```
/\ ~ tmState = _COMMITTED
```

```
/\ timeout = "off"
```

```
/\ timeout' = "on"
```

```
/\ UNCHANGED<<msgs, rmState, tmState, rmPrepared, rmPrecommitted>>
```

Three-Phase Commit

```
RMWhenTimeout(r) ==  
  /\ timeout = "on"  
  /\ rmState[r] \notin {_ABORTED, _COMMITTED}  
  /\ IF rmPrecommitted # {}  
    THEN  
      /\ rmState' = [rmState EXCEPT ![r] = _COMMITTED]  
    ELSE  
      /\ rmState' = [rmState EXCEPT ![r] = _ABORTED]  
  /\ UNCHANGED<<msgs,tmState, rmPrepared, rmPrecommitted, timeout>>
```

Three-Phase Commit

```
RMWhenTimeout(r) ==  
  /\ timeout = "on"  
  /\ rmState[r] \notin { _ABORTED, _COMMITTED }  
  /\ IF rmPrecommitted # {}  
    THEN  
      /\ rmState' = [rmState EXCEPT ![r] = _COMMITTED]  
    ELSE  
      /\ rmState' = [rmState EXCEPT ![r] = _ABORTED]  
  /\ UNCHANGED<<msgs,tmState, rmPrepared, rmPrecommitted, timeout>>
```

```
TMWhenTimeout ==  
  /\ timeout = "on"  
  /\ tmState = _INIT  
  /\ IF rmPrecommitted # {} /\ ~ [ type |-> _ABORT ] \in msgs  
    THEN  
      /\ tmState' = _COMMITTED  
    ELSE  
      /\ tmState' = _ABORTED  
  /\ UNCHANGED<<msgs,rmState, rmPrepared, rmPrecommitted, timeout>>
```

Three-Phase Commit

- Remember:

VARIABLE

rmState, \ * State of each resource manager
tmState, \ * State of the transaction manager
rmPrepared, \ * The set of RMs from which the TM has received prepared messages
rmPrecommitted, \ * The set of RMs from which the TM has received pre-committed messages
msgs, \ * Set of all messages
timeout \ * Timeout flag

Three-Phase Commit

- Remember:

VARIABLE

rmState, * State of each resource manager
tmState, * State of the transaction manager
rmPrepared, * The set of RMs from which the TM has received prepared messages
rmPrecommitted, * The set of RMs from which the TM has received pre-committed messages
msgs, * Set of all messages
timeout * Timeout flag

TypeOK ==

```
/\ rmState \in [ RM -> { _WORKING, _PREPARED, _PRECOMMITTED, _COMMITTED, _ABORTED } ]  
/\ tmState \in { _INIT, _COMMITTED, _ABORTED }  
/\ rmPrepared \subseteqq RM  
/\ rmPrecommitted \subseteqq RM  
/\ msgs \subseteqq Message  
/\ timeout \in { "on", "off" }
```

Three-Phase Commit

Agreement ==

$$\begin{aligned} & \backslash A \text{ } r m 1, r m 2 \backslash \text{in } R M : \sim / \backslash r m S t a t e [r m 1] = _ A B O R T E D \\ & \quad / \backslash r m S t a t e [r m 2] = _ C O M M I T T E D \end{aligned}$$

Validity ==

$$\backslash E r \backslash \text{in } R M : r m S t a t e [r] = _ A B O R T E D \sim \rightarrow t m S t a t e = _ A B O R T E D$$

Fairness ==

$$\begin{aligned} & / \backslash W F _ v a r s (T M C o m m i t) \\ & / \backslash W F _ v a r s (T M A b o r t) \\ & / \backslash W F _ v a r s (T M P r e c o m m i t) \\ & / \backslash \backslash A r \backslash \text{in } R M : W F _ v a r s (T M R c v P r e c o m m i t (r)) \\ & / \backslash \backslash A r \backslash \text{in } R M : W F _ v a r s (T M R c v P r e p a r e (r)) \\ & / \backslash \backslash A r \backslash \text{in } R M : W F _ v a r s (T M R c v A b o r t (r)) \\ & / \backslash \backslash A r \backslash \text{in } R M : W F _ v a r s (T M W h e n T i m e o u t) \\ & / \backslash \backslash A r \backslash \text{in } R M : W F _ v a r s (R M W h e n T i m e o u t (r)) \end{aligned}$$

Three-Phase Commit

Next ==

\\ TMPrecommit

\\ TMCommit

\\ TMAbort

\\ Timeout

\\ TMWhenTimeout

\\ \E r \in RM:

\\ TMRcvPrepare(r)

\\ TMRcvPrecommit(r)

\\ TMRcvAbort(r)

\\ RMPPrepare(r)

\\ RMPrecommit(r)

\\ RMCommit(r)

\\ RMWhenTimeout(r)

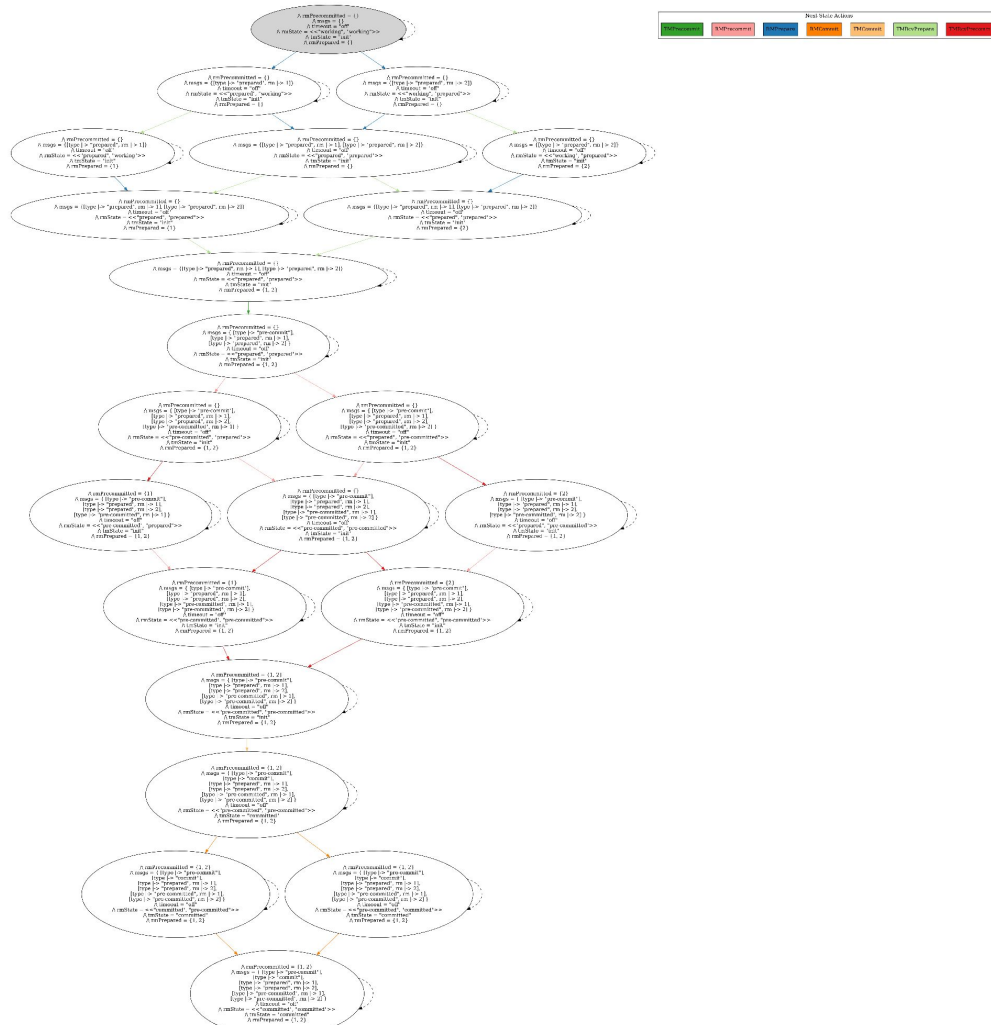
\\ RMChooseToAbort(r)

\\ RMRcvAbortMsg(r)

Our 3PC Specification

TLC reports no errors!

Time	Diameter	States Found	Distinct States	Queue Size
00:00:48	29	256,756	84,111	0
00:00:05	8	42,827	16,907	8,530
00:00:01	0	1	1	1



Evaluation Questions

Conclusion

- TLA+ is a useful tool to verify properties of an algorithm
- Complex systems are hard to specify
- Only validates the provided specification

Conclusion

- TLA+ is a useful tool to verify properties of an algorithm
- Complex systems are hard to specify
- Only validates the provided specification

Conclusion

- TLA+ is a useful tool to verify properties of an algorithm
- Complex systems are hard to specify
- Only validates the provided specification

Conclusion

- TLA+ is a useful tool to verify properties of an algorithm
- Complex systems are hard to specify
- Only validates the provided specification

Future Work

Future Work

Two-Phase Commit

- Allow Transaction manager to crash
- Model message loss / Network partitions

Future Work

Two-Phase Commit

- Allow Transaction manager to crash
- Model message loss / Network partitions

Future Work

Two-Phase Commit

- Allow Transaction manager to crash
- Model message loss / Network partitions

Future Work

Three-Phase Commit

- Verify Validity
- Verify Termination
- Model Network Partitions
- Model Process crashes

Future Work

Three-Phase Commit

- Verify Validity
- Verify Termination
- Model Network Partitions
- Model Process crashes

Future Work

Three-Phase Commit

- Verify Validity
- Verify Termination
- Model Network Partitions
- Model Process crashes

Future Work

Three-Phase Commit

- Verify Validity
- Verify Termination
- Model Network Partitions
- Model Process crashes

Thank you!

Thank you!
Questions?